

Binary Search Trees

Dr. Anirban Ghosh

School of Computing
University of North Florida



Maps

Definition

A **record** is a key-value pair: (k, v)

A **map** is an abstract data type for maintaining a set of records

- No two records can have the same key
- However, two records can have same values though
- Association of keys to values define a **mapping**: $f(\text{key}) = \text{value}$

Examples of maps

- UNF maintains a map of (N#, student information) records
- A social media company maintains a map of (email address, user account information) records
- An assembler maintains a symbol table (a map) of (opcode, hex) records
- A text-editor maintains a map of (color, RGB representation) records

How to implement a map?

Common map operations

- **Insert** a record (k, v)
- **Retrieve** a record having key k
- **Delete** a record having key k

Approach 1: maintain a sorted list of records

- **Insertion.** will take $O(\log n)$ time for figuring out the correct spot for the incoming record using binary search; then $O(n)$ time for shifting items to the right to accommodate the new record; total time taken is $O(\log n) + O(n) = O(n)$
- **Retrieval.** will take $O(\log n)$ time using binary search
- **Deletion.** will take $O(\log n)$ time to locate it using a binary search; then then $O(n)$ time for left shifting items to fill the empty spot; total time taken $O(\log n) + O(n) = O(n)$

How to implement a map?

Common map operations

- **Insert** a record (k, v)
- **Retrieve** a record having key k
- **Delete** a record having key k

Approach 2: maintain an unsorted list of records

- **Insertion.** will take $O(1)$ time (add the new record at the end)
- **Retrieval.** will take $O(n)$ time using a linear search; may need to search the whole list in the worst-case
- **Deletion.** will take $O(n)$ time to locate it using a linear search; then $O(n)$ time for left shifting items to fill the empty spot; total time taken $O(n) + O(n) = O(n)$

Common map operations

- **Insert** a record (k, v)
- **Retrieve** a record having key k
- **Delete** a record having key k


- ☞ To accomplish the above three tasks in $O(\log n)$ time each
- ☞ **Balanced** binary search trees is the solution; stay tuned ...

The map ADT


```
public interface MapADT<K,V> {  
    boolean put(K k, V v); // adds a new record with key k and value v  
    V remove(K k); // removes the record having key k  
    V get(K k); // return the value part of the record whose key is k  
    V updateValue(K k, V v); // updates the value part of the record whose key is k with a new value v  
    int size(); // returns the number of records stored in the map  
    void clear(); //Removes all records from the map  
}
```

What is a Binary Search Tree?

- It is a **binary tree** where every node contains a **<key, value>** pair (a **record**); keys must be **comparable** but the values don't need to be
- Moreover, for every node p in the tree, the following 2 properties hold
 - 1 Keys stored in the left subtree of p are $<$ the key stored at p
 - 2 Keys stored in the right subtree of p are $>$ the key stored at p



What is a Binary Search Tree?

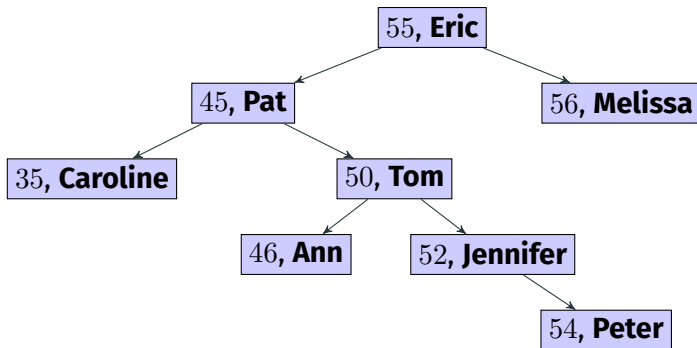


Use

BSTs can be used to implement **maps** and are commonly used for fast searching (typically need far less comparisons than lists)

An important property of BSTs

An inorder traversal of a BST always gives the sorted sequence based on the keys



Inorder traversal

35, **Caroline**; 45, **Pat**; 46, **Ann**; 50, **Tom**; 52, **Jennifer**; 54, **Peter**; 55, **Eric**; 56, **Melissa**


Let's say you need to look for the record that has the key k ; how will you do this?

Algorithm

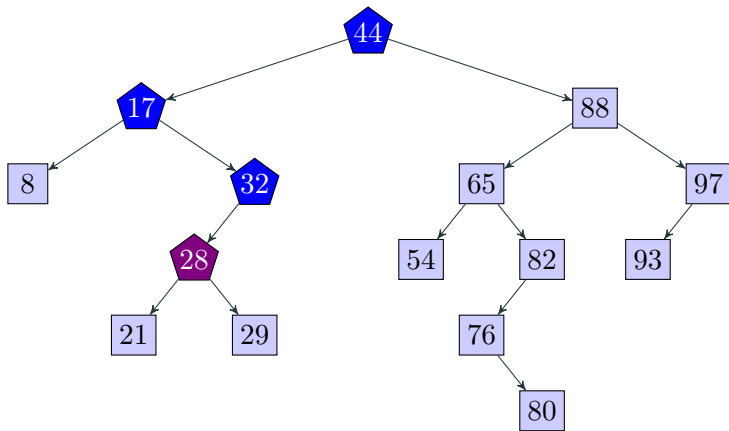
- Start at the root
- If the root's key is k , then search is successful
- If $k < \text{root's key}$, search recursively (or iteratively) in the left subtree of the root
- Otherwise, search recursively (or iteratively) in the right subtree of the root
- If we have reached a null link, no record exists in the tree with key k

To save space in the figures, we will write only the record keys inside the nodes and avoid the corresponding values

Example




Example: search for 28



Found!

Example: search for 68



Not found; a null link is reached (the left link of 76)

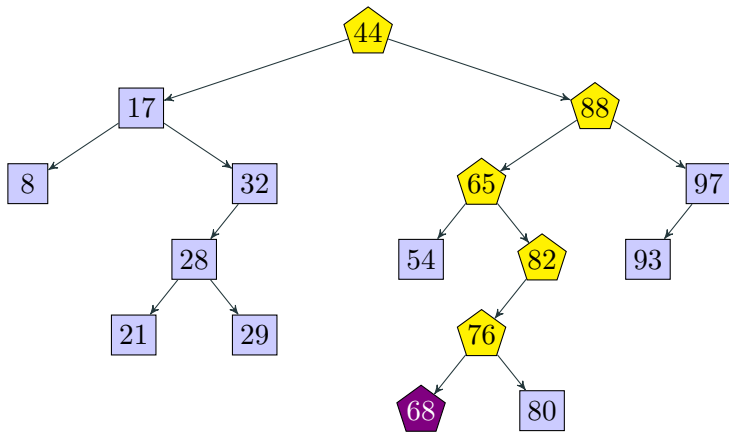
$O(h)$, where h is the height of the BST under consideration, since, for searching, we need to traverse a path whose length is h in the worst-case

How to insert a record into a BST having key k ?

Algorithm

- Start at the root
- If the root is null, replace empty tree with a new tree with the new record as the root, and signal **SUCCESS**
- If k equals root's key, signal **FAILURE** since a record with key k already exists
- If $k <$ root's key, insert recursively (or, iteratively) in the left subtree of the root
- Otherwise, insert recursively (or, iteratively) in the right subtree of the root

Example: insert 68



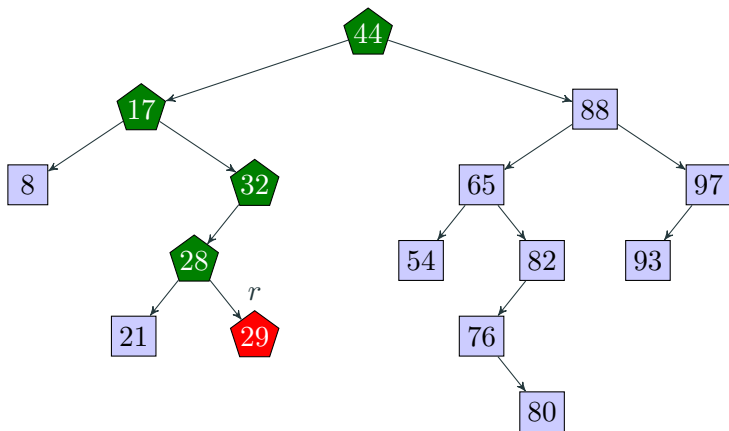
$O(h)$, where h is the height of the BST under consideration, since, for inserting a new record, we need to traverse a path whose length is h in the worst-case

How to delete the record from a BST having key k ?

Idea

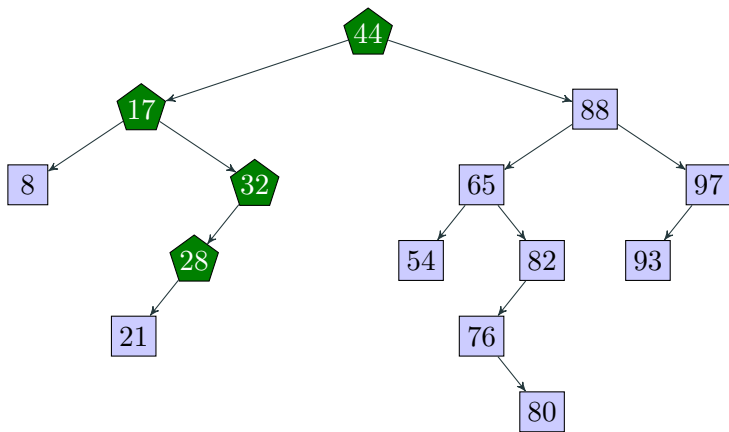
- Traverse through the tree to find the record having key k
- If the record cannot be found, no action is needed
- Now assume that we have found the record that has key k at node r
 - 1 r is a leaf node (no child)
 - 2 r has exactly one child (either left or right)
 - 3 r has two children (both left and right)

Case 1: r is a leaf node, delete 29

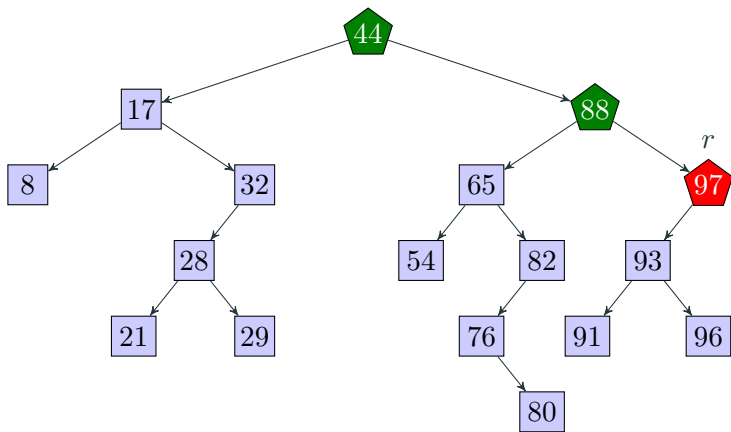


Easy! just delete it right away

Case 1: r is a leaf node, delete 29

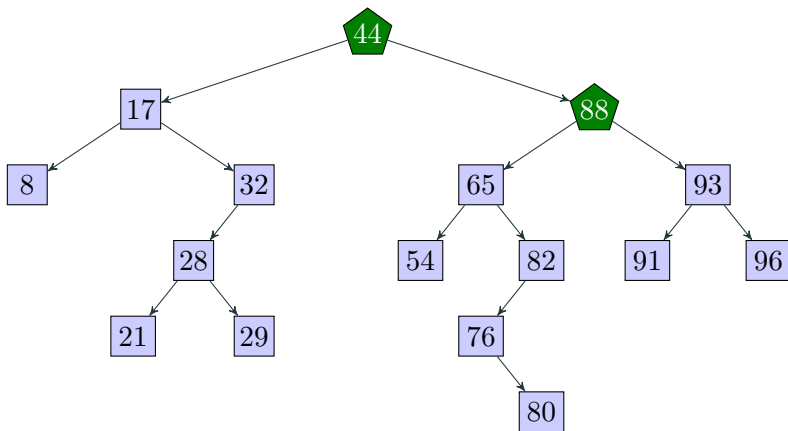


Case 2: r has one child, delete 97

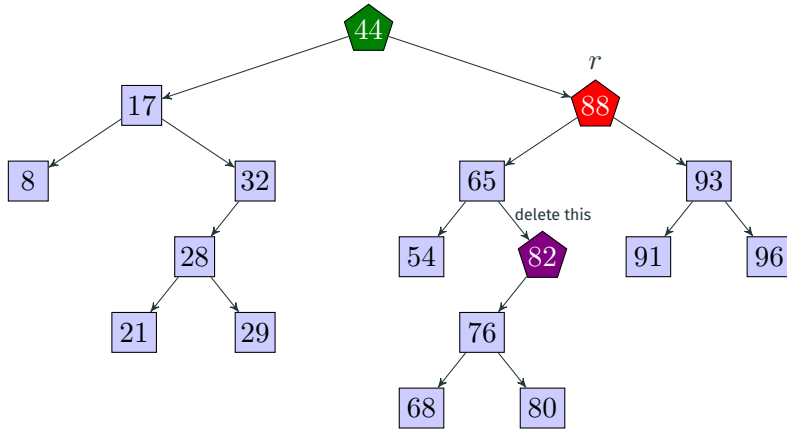


Make the parent of r point to the only child of r instead of r

Case 2: r has one child, delete 97




Case 3: r has two children, delete 88




Replace the record at r with the record at the inorder predecessor p of n ; then delete p using Case 1 or 2, depending on the number of children of p . Note that the inorder predecessor is either a leaf node or has a left child only (no right child).

Case 3: r has two children, delete 88




Replace the record at r with the record at the inorder predecessor p of n ; then delete p using Case 1 or 2, depending on the number of children of p . Note that the inorder predecessor is either a leaf node or has a left child only (no right child).

Case 3: r has two children, delete 88



Replace the record at r with the record at the inorder predecessor p of n ; then delete p using Case 1 or 2, depending on the number of children of p . Note that the inorder predecessor is either a leaf node or has a left child only (no right child).

Finding the inorder predecessor in Case 3



The inorder predecessor p of node r is either a leaf node (left figure) or has a left child but no right child (right figure). So, to find p , go left of r and then keep on moving right as long as possible. It takes $O(h)$ time to find p , where h is the height of the tree. The yellow triangles represent subtrees (some of them could be empty).

Time complexity of deletion

- It takes $O(h)$ time for locating the record to be deleted, where h is the height of the BST
- Case 1 takes $O(1)$ time to delete a record
- Case 2 takes $O(1)$ time to delete a record
- Case 3 takes $O(h)$ time to delete a record since we need to locate the inorder predecessor p of node r in $O(h)$ time and then delete p in $O(1)$ time using Case 1 or 2
- So, a single deletion takes $O(h)$ time in the worst-case

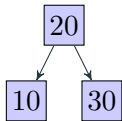
A built-in record class in Java

```
java.util.AbstractMap.SimpleEntry<K,V>
```

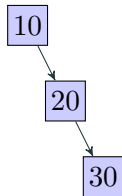
Reference. <https://docs.oracle.com/en/java/javase/17/docs/api/java.base/java/util/AbstractMap.SimpleEntry.html>

See the class `TreeMapBST`

BSTs are not unique



Insertion sequence: 20, 10, 30



Insertion sequence: 10, 20, 30



***Their structures really depend
on the insertion sequence of records***

An application

Problem


Given a text, find out the unique words in it along with their counts. We also need to output the distinct words with their count.

A solution

- Use a BST T where the key-type is `String` and the value-type is `Integer`. This means every node will store a word from the text and its count in the same text.
- For every word w in the text, first check if a node exists in T , where the stored key is w
 - If such a node does not exist, insert a new node in T with key w and value 1
 - If such a node exists in T , increment the stored value (essentially a counter) by 1

See the class [WordCounter](#)

Worst case scenario: skewed binary trees



In this case, $k_1 < k_2 < k_3 < \dots < k_n$ and the tree is right-skewed
Note that when $k_n < k_{n-1} < \dots < k_1$, the tree will be left-skewed (try!)

So, in the worst-case, $h = n - 1 = O(n)$

Therefore, searching, insertion, deletion take $O(h) = O(n)$ time each!

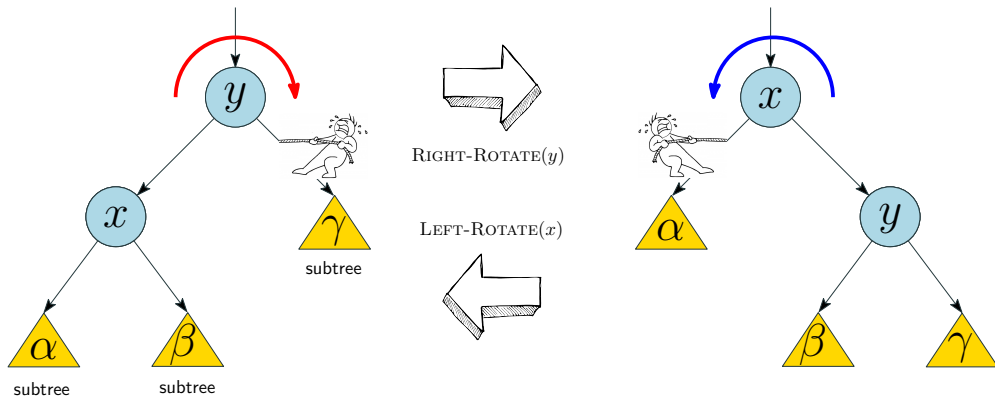
This is as bad as using singly linked-lists!

Do something so that h remains under control (logarithmic)
We aim for $h = O(\log n)$

A solution. use **Red-Black** trees

Red-Black trees are never skewed or close to being skewed unlike the plain binary search trees we just talked about 👍

Rotation on BSTs

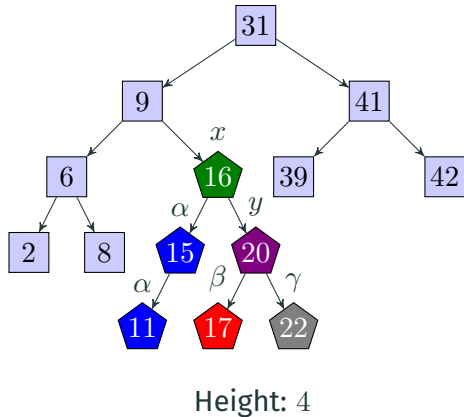
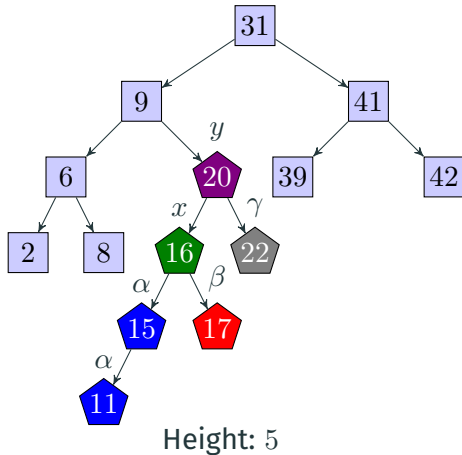


Rotations help in **reducing** height of BSTs; this means faster operations on BSTs

A single rotation can be done in $O(1)$ time

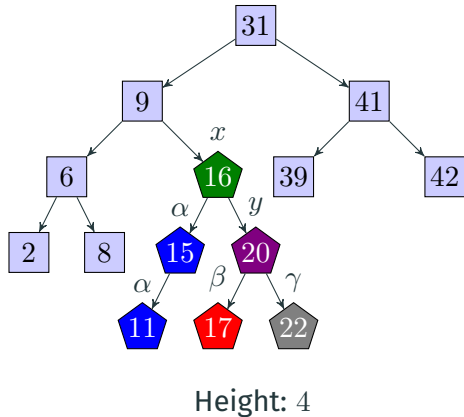
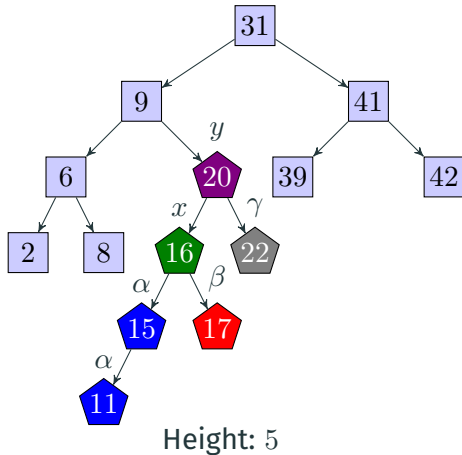
☞ Some of the three subtrees (α, β, γ) can be empty (devoid of nodes).

An example of rotation



- ☞ If a **right** rotation is performed at node y in the left tree, we get the tree on the right.
- ☞ If a **left** rotation is performed at node x in the right tree, we get back the tree on the left.

An example of rotation



Note that rotations never alter the inorder traversal sequences. It also means that after a rotation, the resulting binary tree is still a binary search tree.

Code for rotation

Right rotation at node y

```
private void rightRotateAt(Node<K,V> y) {
    Node<K,V> x = y.left;
    y.left = x.right;

    if( x.right != null )
        x.right.parent = y;

    x.parent = y.parent;

    if( y.parent == null )
        root = x;
    else if( y == y.parent.right )
        y.parent.right = x;
    else
        y.parent.left = x;

    x.right = y;
    y.parent = x;
}
```

Left rotation at node x

```
private void leftRotateAt(Node<K,V> x) {
    Node<K,V> y = x.right;
    x.right = y.left;

    if( y.left != null )
        y.left.parent = x;

    y.parent = x.parent;

    if( x.parent == null )
        root = y;
    else if( x == x.parent.left )
        x.parent.left = y;
    else
        x.parent.right = y;

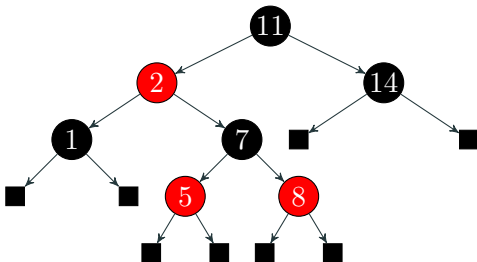
    y.left = x;
    x.parent = y;
}
```

A single rotation (left/right) takes $O(1)$ time

Definition

A **red-black** tree is a self-balancing BST (height never gets too bad) that has the following properties:

- 1 Every node is either **red** or **black**
- 2 The root is **black**
- 3 If a node is **red**, then both its children are **black**. The **null** links (shown using black squares in the figure) of the leaves are **black**.
- 4 The number of **black** nodes in any path from the root to a leaf is the same



How to maintain the colors?

Since the colors are bichromatic, a boolean variable `color` is enough to specify the color of a node. Set `color` to `false`, when the node is painted **RED**, or else, set it to `true` to specify the color is **BLACK**.

☞ There is no magic behind these two colors; feel free to use any two colors

A node class implementing for RB-tree

```
public class TreeMapRBTree<K extends Comparable<K>, V> implements MapADT<K,V>, Iterable< SimpleEntry<K,V> > {  
    private static final boolean RED = false, BLACK = true;  
  
    private static class Node<K, V> {  
        final private K key;  
        private V val;  
  
        private Node<K, V> left, right, parent;  
  
        private boolean color;  
  
        public Node(K k, V v) {  
            key = k; val = v;  
            left = right = parent = null;  
            color = RED;  
        }  
  
        public String toString() {  
            String colorString = (color == RED)? "RED" : "BLACK";  
            if (val != null) return "<" + key.toString() + ", " + val.toString() + ", " + colorString + ">" ;  
            else return key.toString();  
        }  
  
        // other variables and methods  
    }  
}
```


Height of a red-black tree

- A RB-tree containing n records has height $\leq 2\log_2(n+1) = O(\log n)$
- **Intuition.** By constraining the node colors on any simple path from the root to a leaf (property 4), it can be ensured that no such path is more than twice as long as any other so that the tree's height is always logarithmic.

Implication

When $n = 1,000,000$, $h \leq 2\log_2(n+1) < 40$.

To search for a record, at most 41 comparisons are required in this case.

In contrast, when a plain binary search tree is used, $h \leq 999,999$.

For searching for a record, at most 1,000,000 comparisons are required in this case.

This clearly demonstrates the advantage of using RB-trees over plain BSTs and also shows that RB-trees can never be skewed.

👉 *Searching, insertion, and deletion in RB-trees take $O(h) = O(\log n)$ time each.*

The three primary operations

- ❶ **Search.** same as the search operation for plain BSTs; takes $O(h) = O(\log n)$ time (note that RB-trees are also BSTs, so the same search algorithm works here too!)
- ❷ **Insertion.** we **will** discuss this; takes $O(\log n)$ time
- ❸ **Deletion.** we **won't** discuss this; takes $O(\log n)$ time

👉 The TreeMap class in Java implements RB-Tree

<https://docs.oracle.com/en/java/javase/17/docs/api/java.base/java/util/TreeMap.html>

Inserting a new node z into a RB-tree

- 1 Color z **red** and insert it as you would in a plain BST
- 2 If necessary, start fixing the tree (using rotations and recoloring) as long as you see z 's parent is **red** (z may change as we climb up the tree); see the cases next

```
while( z.parent != null && z.parent.color == RED ){  
    // deal with the cases inside this loop  
}
```

- 3 At the end, color the root using **black**

```
root.color = BLACK;
```

Terminologies

Uncle of a node

The uncle of a node is the sibling of its parent. In some cases, it could be a null link if there is no such sibling node.

```
private boolean isLeftChild( Node<K,V> node ) {  
    return node.parent != null && node.parent.left == node;  
}  
  
private Node<K,V> uncle( Node<K,V> node ){  
    if( node == null ) return null;  
    return ( isLeftChild(node.parent) ) ? node.parent.parent.right : node.parent.parent.left;  
}
```

Grandparent of a node

The grandparent of a node is the parent of its parent. In some cases, it could be a null link if there is no such grandparent node.

```
private Node<K,V> grandParent( Node<K,V> node ){  
    return node.parent.parent;  
}
```

The six cases of insertion

- (Case 1a) z 's uncle y is **RED** and z is a right child
- (Case 1b) z 's uncle y is **RED** and z is a left child
- (Case 2a) z is a right child, z 's parent is a left child, and z 's uncle y is **BLACK**
- (Case 2b) z is a left child, z 's parent is a right child, and z 's uncle y is **BLACK**
- (Case 3a) z is a left child, z 's parent is left child, and z 's uncle y is **BLACK**
- (Case 3b) z is a right child, z 's parent is a right child, and z 's uncle y is **BLACK**

Things to note

- 1 The first node inserted into an empty red-black tree is always painted **BLACK**
- 2 Since the null links are considered to be **BLACK**, make sure that you consider their existence and color while finding the uncle of nodes
- 3 In each of the six cases, the parent of z is always **RED** (if it is **BLACK**, no fixing would be required)

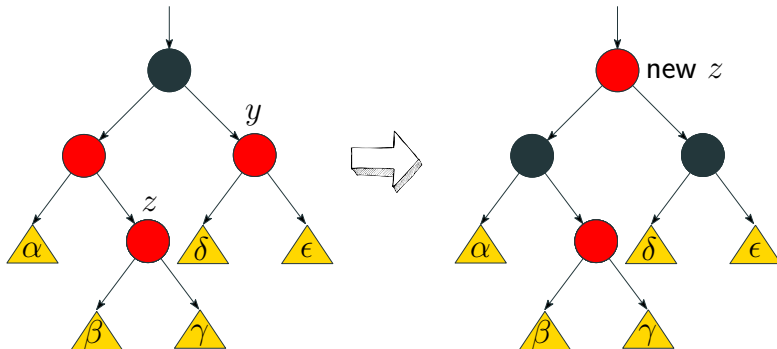
Case 1



(Case 1a) z 's uncle y is **RED** and z is a right child

(Case 1b) z 's uncle y is **RED** and z is a left child

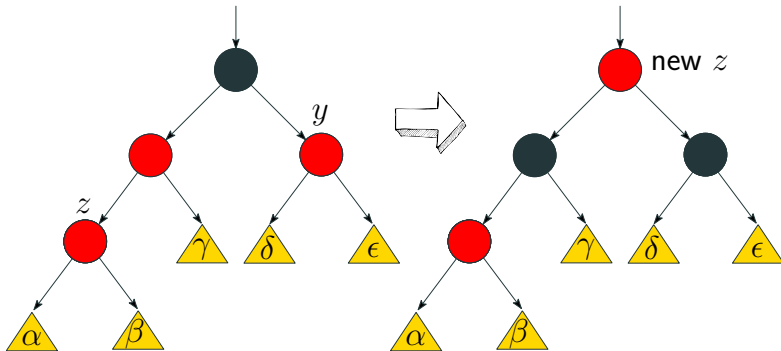
Case 1a



z 's uncle y is **red** and z is a right child; recoloring is needed but no rotation; takes $O(1)$ time; **continue fixing up** using the new z node

- `z.parent.color = BLACK; y.color = BLACK; grandParent(z).color = RED;`
- `z = grandParent(z); // continue fixing up using the new 'z'`

Case 1b



z 's uncle y is **red** and z is a left child; recoloring is needed but no rotation; takes $O(1)$ time; **continue fixing up** using the new z node

- `z.parent.color = BLACK; y.color = BLACK; grandParent(z).color = RED;`
- `z = grandParent(z); // continue fixing up using the new 'z'`

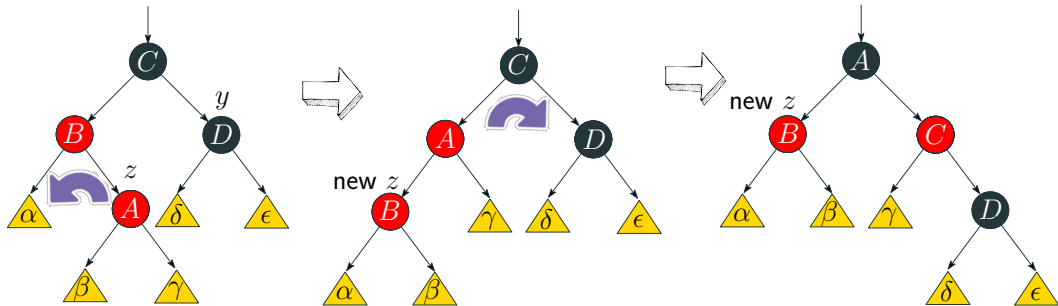
Case 2



(Case 2a) z is a right child, z 's parent is a left child, and z 's uncle y is **BLACK**

(Case 2b) z is a left child, z 's parent is a right child, and z 's uncle y is **BLACK**

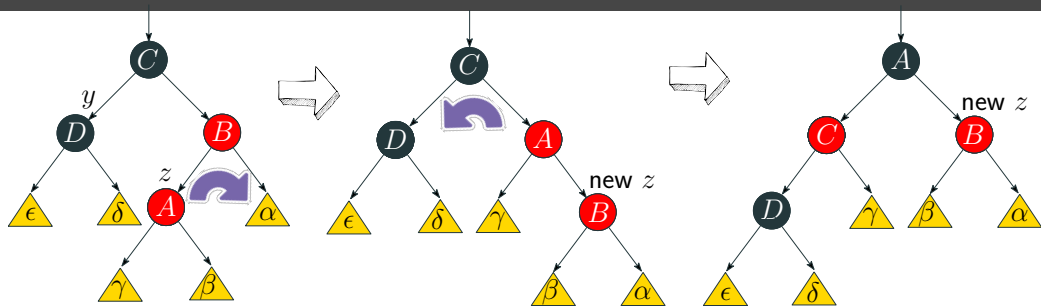
Case 2a



z is a right child, z 's parent is a left child, and z 's uncle y is **black**; recolorings + 2 rotations are needed; the fixing-up process **terminates** since z 's parent is **black** after the two rotations; takes $O(1)$ time

- `z = z.parent; leftRotateAt(z);`
- `z.parent.color = BLACK; grandParent(z).color = RED; rightRotateAt(grandParent(z));`

Case 2b



z is a left child, z 's parent is a right child, and z 's uncle y is **black**; recolorings + 2 rotations are needed; the fixing-up process **terminates** since z 's parent is **black** after the two rotations; takes $O(1)$ time

- `z = z.parent; rightRotateAt(z);`
- `z.parent.color = BLACK; grandParent(z).color = RED; leftRotateAt(grandParent(z));`

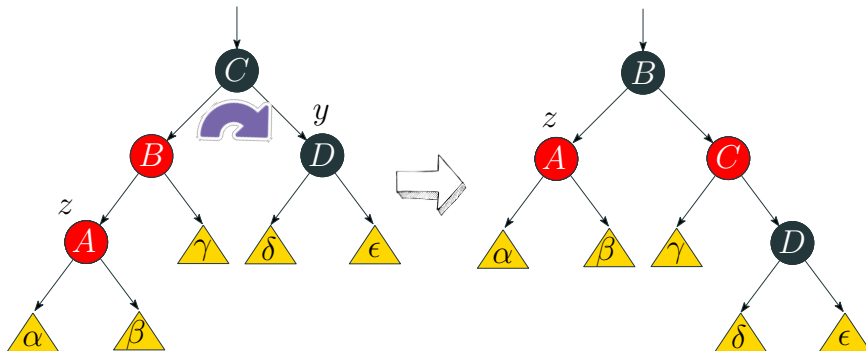
Case 3



(Case 3a) z is a left child, z 's parent is left child, and z 's uncle y is **BLACK**

(Case 3b) z is a right child, z 's parent is a right child, and z 's uncle y is **BLACK**

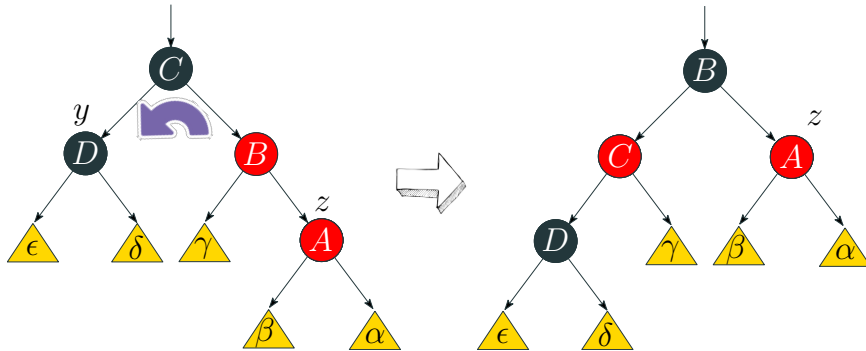
Case 3a



z is a left child, z 's parent is left child, and z 's uncle y is **black**; recolorings + 1 right rotation are needed; the fixing-up process **terminates** since z 's parent is **black** after the rotation; takes $O(1)$ time

```
• parent(z).color = BLACK; grandParent(z).color = RED; rightRotateAt(grandParent(z));
```

Case 3b



z is a right child, z 's parent is a right child, and z 's uncle y is **black**; recolorings + 1 left rotation are needed; the fixing-up process **terminates** since z 's parent is **black** after the rotation; takes $O(1)$ time

```
• parent(z).color = BLACK; grandParent(z).color = RED; leftRotateAt(grandParent(z));
```

The second rotation of Case 2a is exactly an execution of Case 3a
Similarly, the second rotation of Case 2b is exactly an execution of Case 3b

See the class `TreeMapRBTree`

In-browser visualization

41, 38, 31, 12, 19, 8

The value parts of the above records are ignored in this example

ITEM	ACTION
41	The only node, just color it black
38	Insert it to the left of 41; its parent is black , so, no action is needed
31	Case 3a; 31's uncle (a null reference) is black ; right rotate at 41
12	Case 1b; a simple recoloring is enough
19	Case 2a; two rotations are needed
8	Case 1b; a simple recoloring is enough

<https://www.cs.csubak.edu/~msarr/visualizations/RedBlack.html>

Observations

- Time taken for inserting a new node is $O(h) = O(\log n)$; after this fix-up may be needed
- During fix-up, we climb up the tree using Case 1, which only recolors but never rotates
- If we ever use Case 2 (uses 2 rotations) or 3 (uses 1 rotation), we are done (fix-up process terminates)!
- This means at every insertion of a new item, at most 2 rotations are needed
- At most h executions of Case 1 are needed plus 1 execution of Case 2/3, each taking $O(1)$ time
- So, the total time taken for fix-up is $(h + 1) \times O(1) = O(\log n) \times O(1) = O(\log n)$, since, for RB-trees, $h = O(\log n)$
- Total time taken for one insertion equals time taken for inserting a new node plus total time taken for fix-up $= O(\log n) + O(\log n) = 2 \times O(\log n) = O(\log n)$
- Time complexity of searching is $O(h) = O(\log n)$, since, for RB-trees, $h = O(\log n)$
- Deletion also takes $O(h) = O(\log n)$ time but we are not going to discuss it

Plain BSTs vs RB-Trees: tree creation time (using n insertions)

n	RB-tree	Plain BST
10	1	1
100	1	1
1000	3	13
10000	7	214
100000	20	17249

When the input records were already **sorted** in ascending order, RB-trees could easily beat plain BSTs since the heights of the plain BSTs were exactly $n - 1$ everywhere (much worse than the logarithmic heights of RB-trees).

n	RB-tree	Plain BST
10	1	1
100	1	1
1000	3	2
10000	7	8
100000	72	49

When the input records were in **random** order, plain BSTs performed quite as fast as RB-trees since the heights of plain BSTs were not $n - 1$ or even close (in fact, they were almost logarithmic like RB-trees).

👉 Times are reported in milliseconds

Are plain BSTs completely useless?

- Plain BSTs perform **terribly** when the inputs are sorted (or, almost sorted) in ascending or descending order
- But, BSTs are found to perform great on randomly ordered inputs
- In those cases, h is found to be much less than $n - 1$ and is almost logarithmic
- Consequently, we get to see fast searching, insertion, and deletion times
- **Example.** when $n = 5000$, heights of plain BSTs are around 30 if the input is randomly ordered. Note that this is much less than 4999 (worst case height)

<https://opensa-server.cs.vt.edu/ODSA/Books/Everything/html/BST.html>